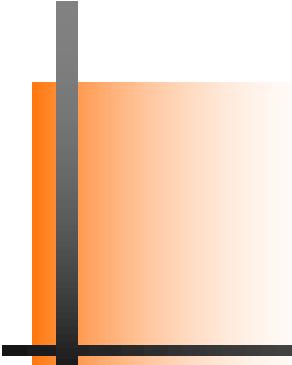


**IN2140:**  
**Introduction to Operating Systems and Data Communication**



# **Operating Systems: Introduction**

---

Pål Halvorsen

# Overview

---

- Basic execution environment – an Intel example  
(why do you need an operating system (OS), or need to know anything about it)
- What is an OS?
- OS components and services  
(extended in later lectures)
- Interrupts and system calls
- Booting, protection, kernel organization



# Hardware

---

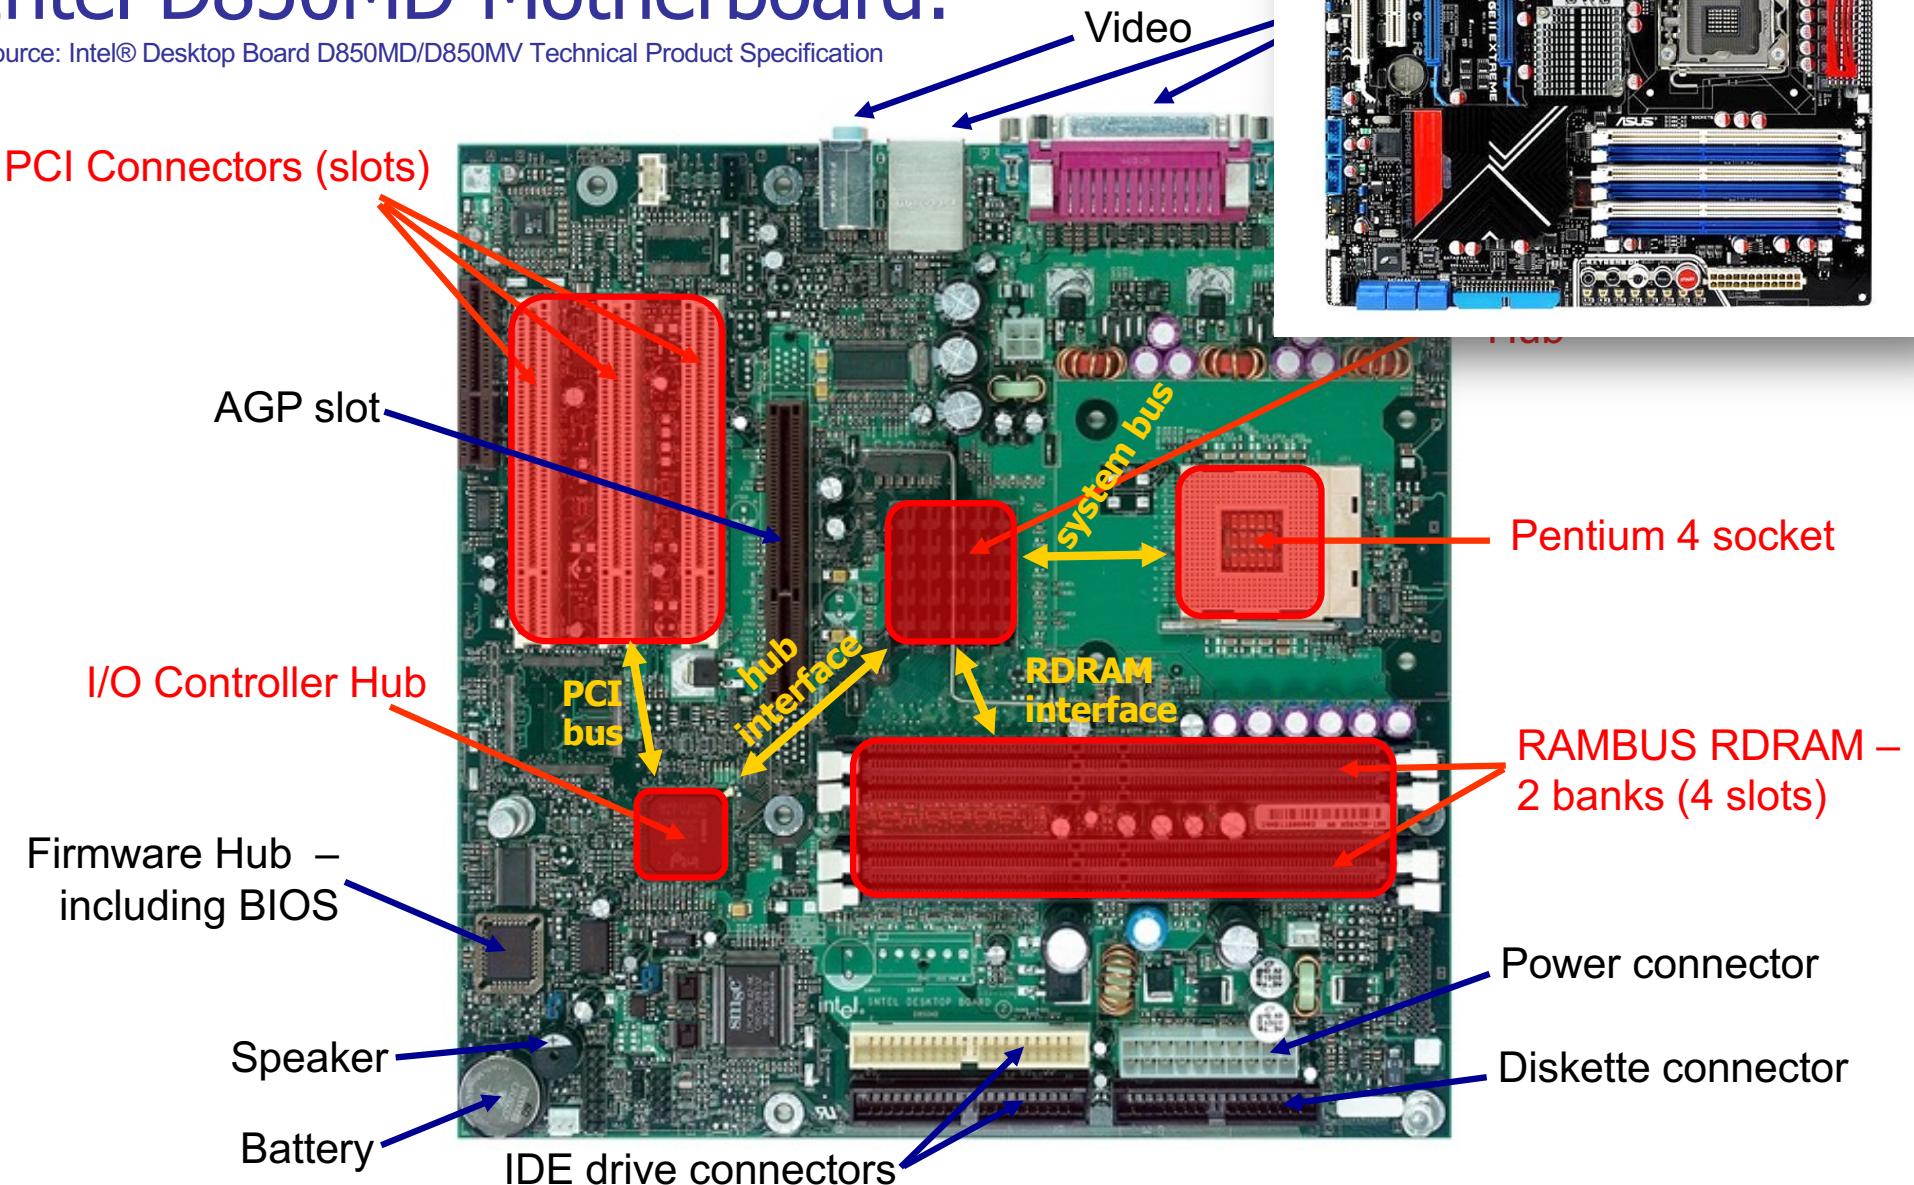
- **Central Processing Units (CPUs)**
- **Memory**  
(cache(s), RAM, ROM, Flash, ...)
- **I/O Devices**  
(network cards, disks, CD, keyboard, mouse, ...)
- **Links**  
(interconnects, busses, ...)



# An easy, old example: Intel Hub Architecture

## Intel D850MD Motherboard:

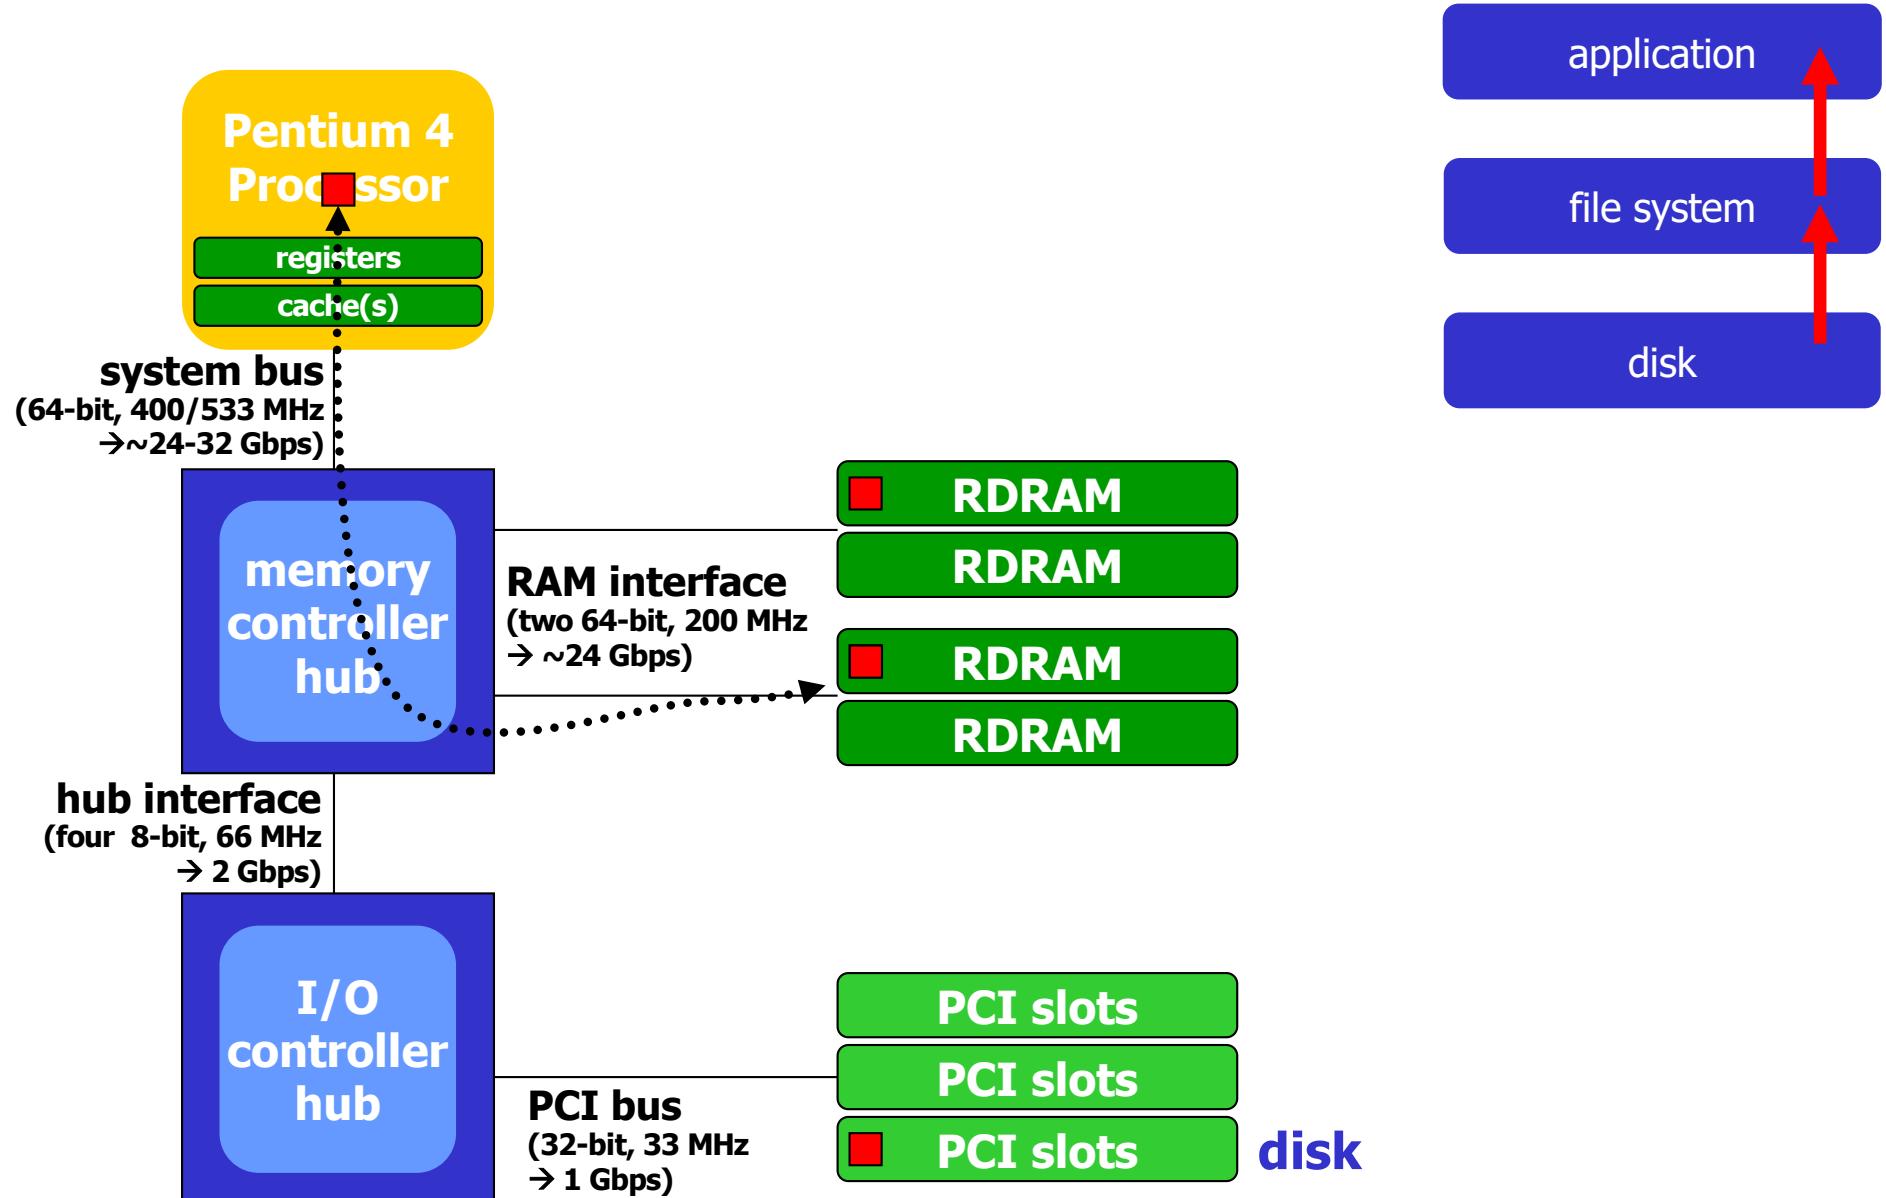
Source: Intel® Desktop Board D850MD/D850MV Technical Product Specification



i7 motherboard



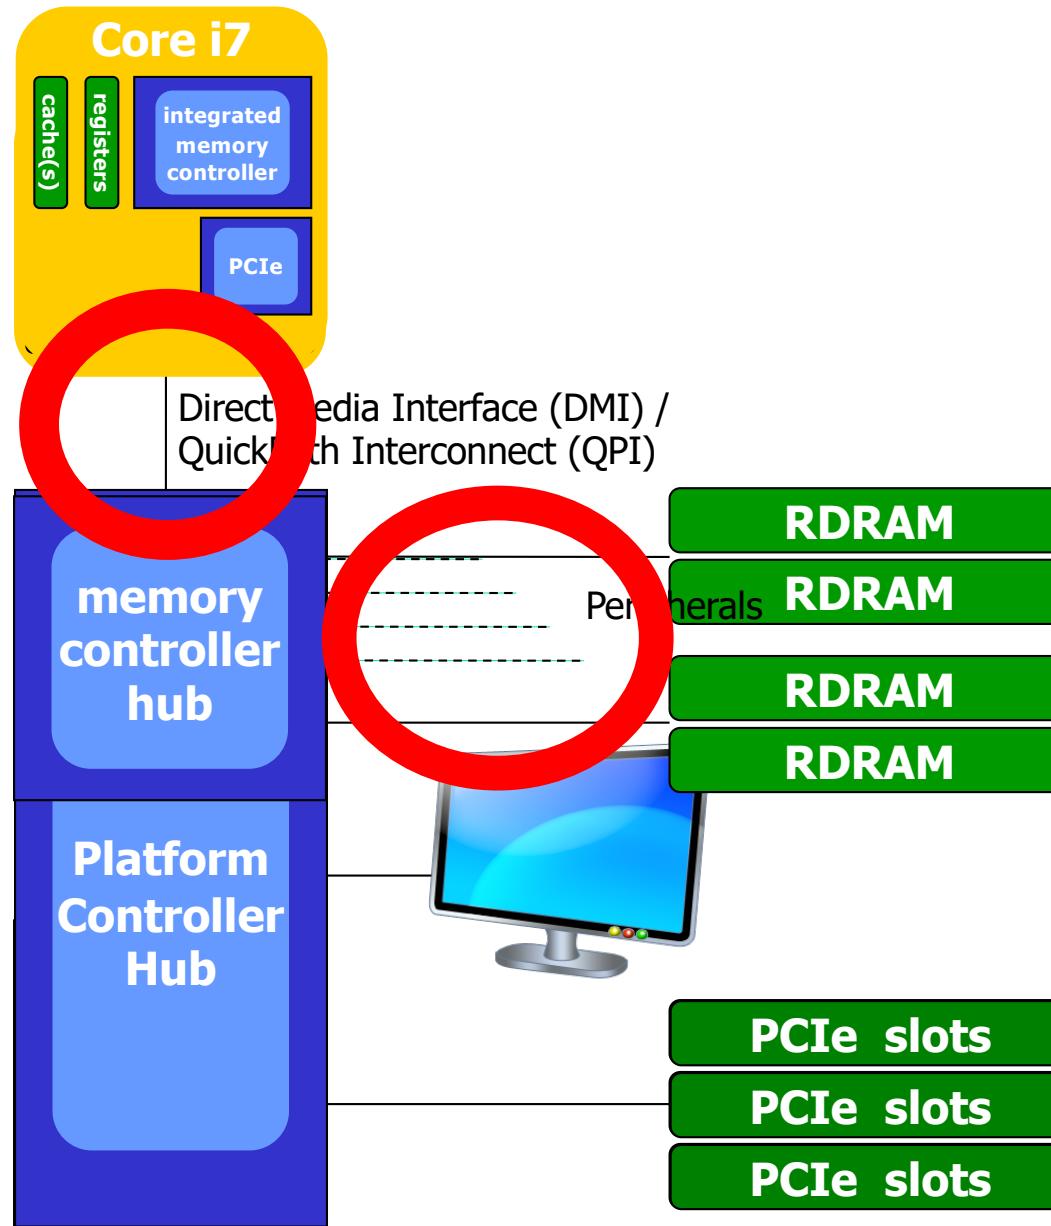
# An easy, old example: Intel Hub Architecture



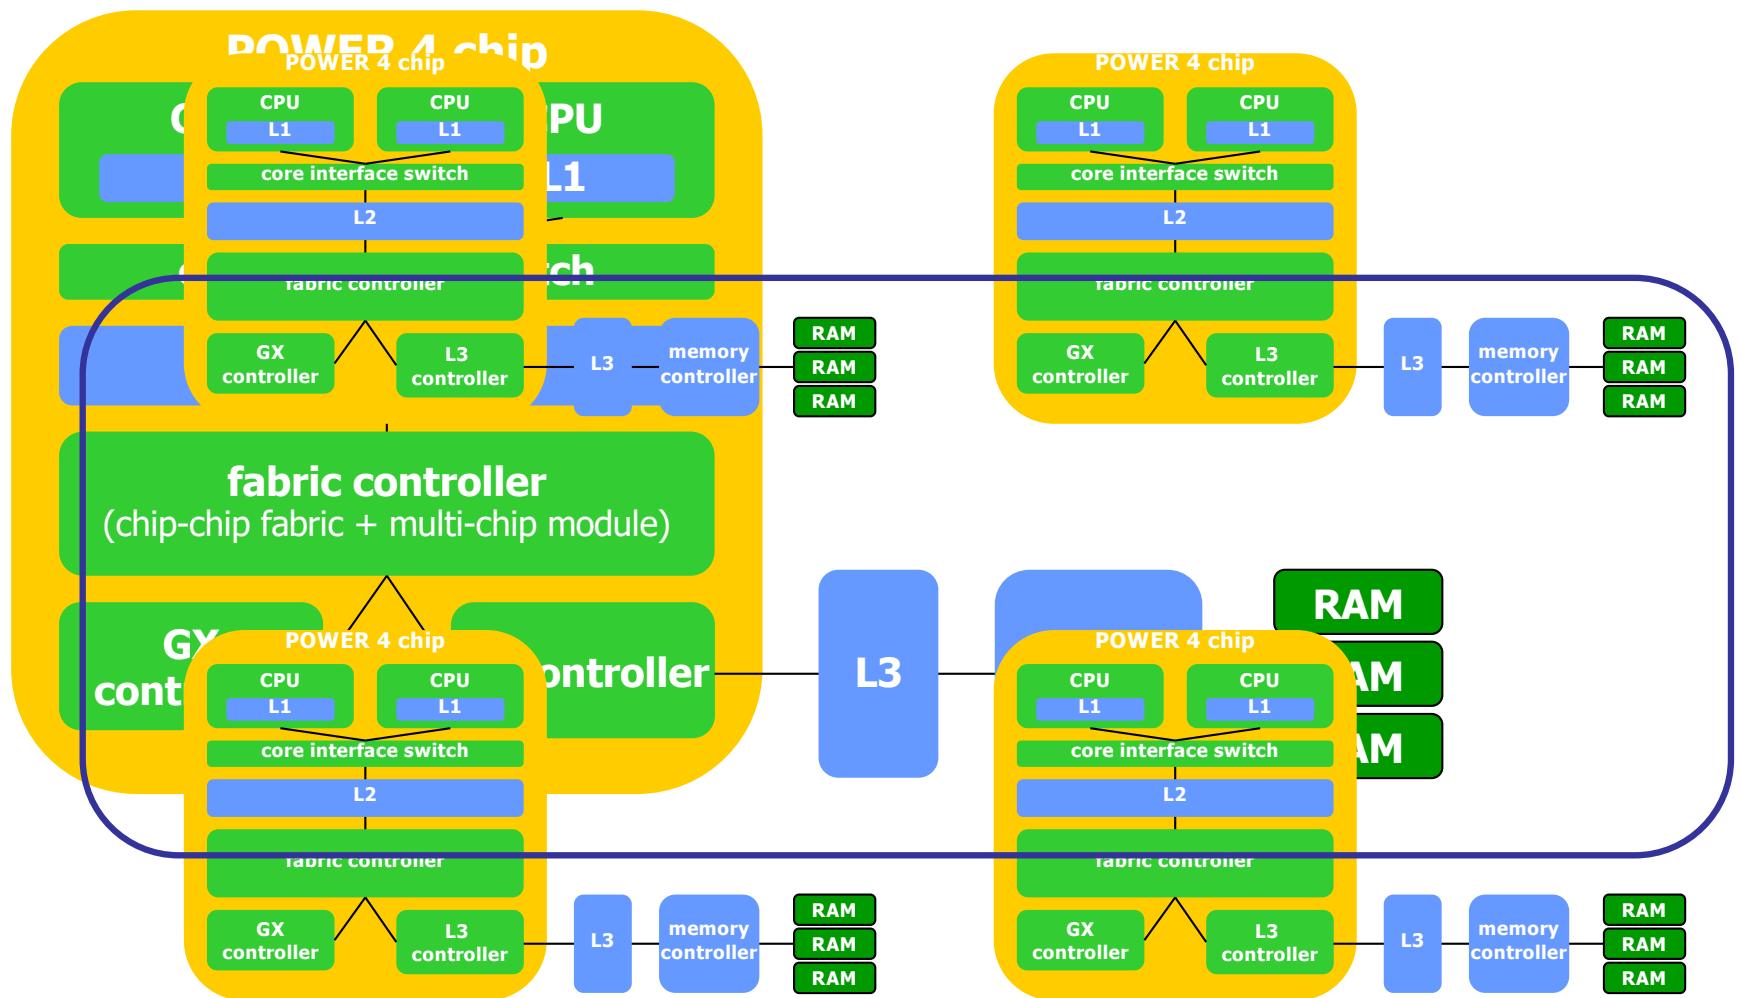
A slightly newer example:

# Intel Platform Controller Hub Architecture

Sandy Bridge, Core i7



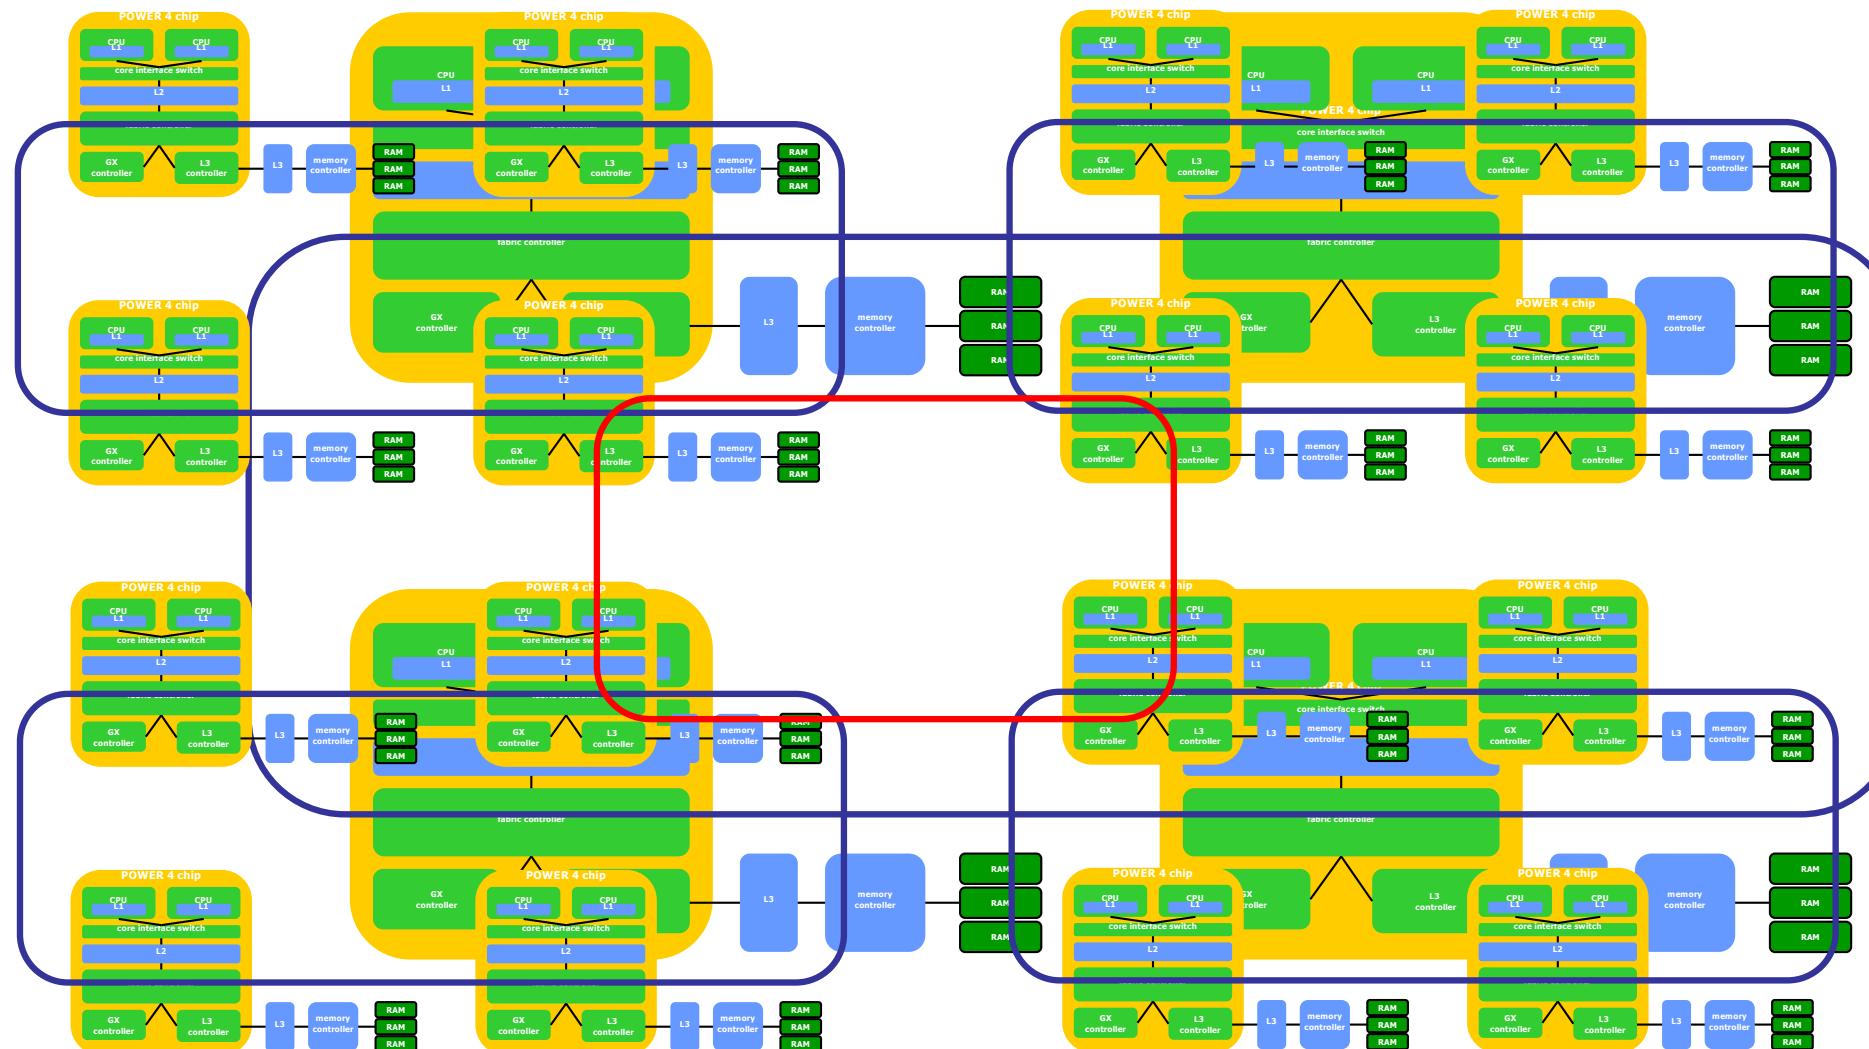
# Example: IBM POWER 4



**Multichip modules** in fabric controller can connect  
4 chips into a 4 chip, 2-way SMP → 8-way MP



# Example: IBM POWER 4



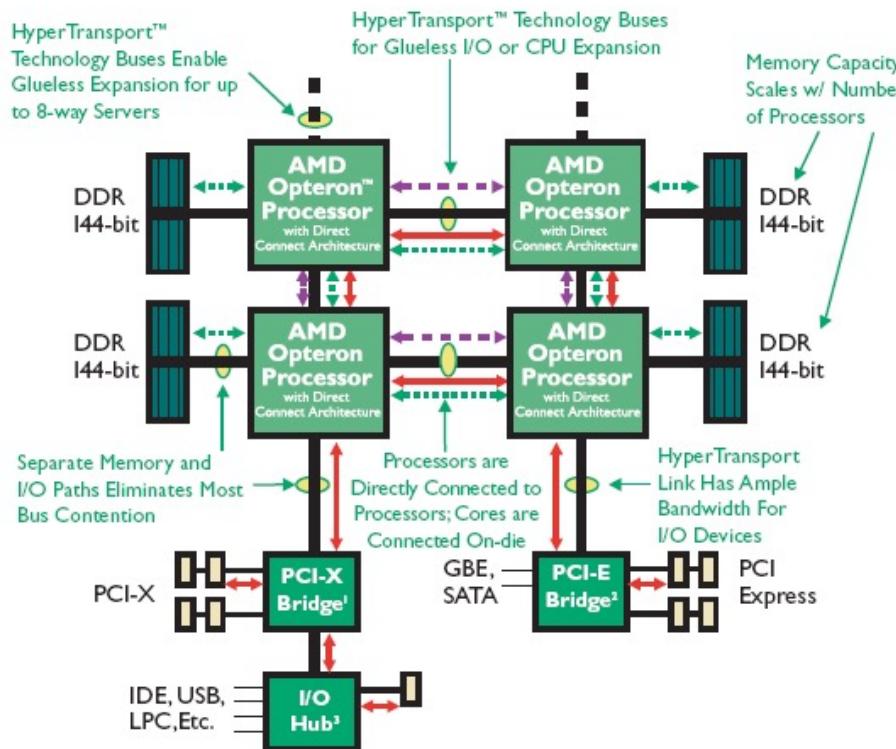
**Chip-chip fabric** in fabric controller can connect 4 multi-chips into a  
4x4 chip, 2-way SMP → 32-way MP



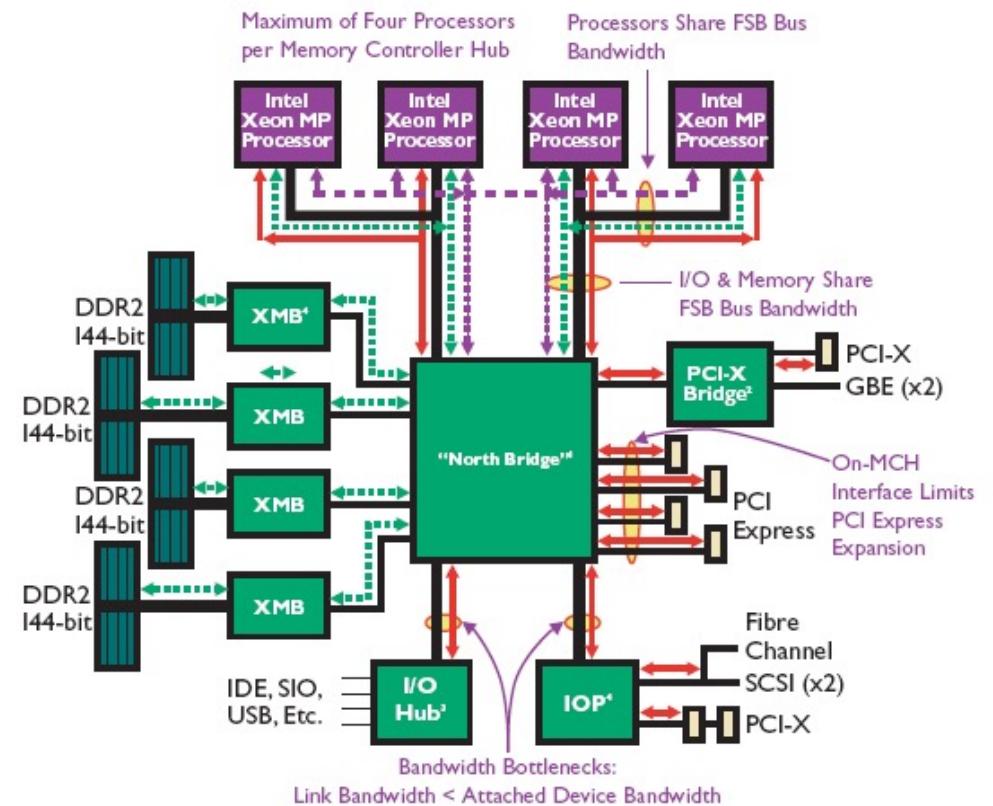
Other examples:

# AMD Opteron & Intel Xeon

## AMD Opteron™ Processor-based 4P Server

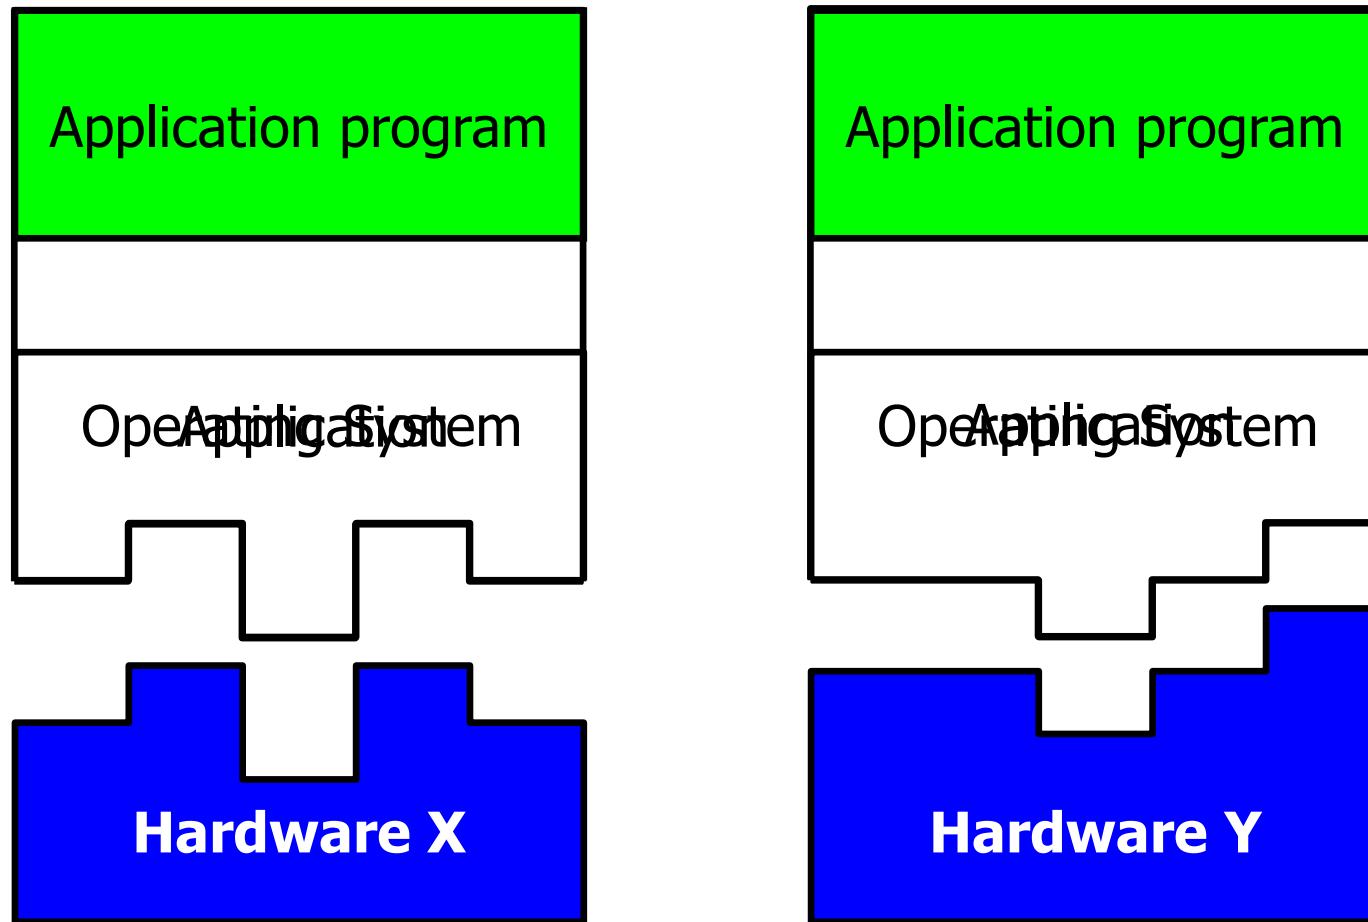


## Intel Xeon MP Processor-based 4P



- ☞ Different hardware may have different bottlenecks  
==> nice to have an **operating system** to control the HW?

# Different Hardware



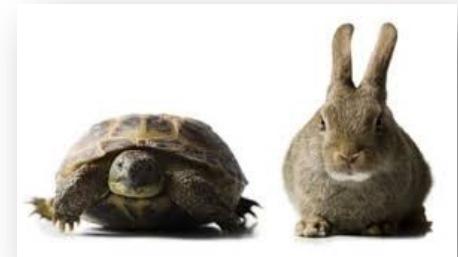
# Intel 32-bit Architecture (IA32): Basic Execution Environment



## 16-bit vs. 32-bit vs. 64 bit vs. ... ?

A big difference between 32-bit processors and 64-bit processors is the size/width of the hardware. This affects the amount of data processed per second – the speed at which they can complete tasks. Otherwise, a lot of similarities...

... using 32-bit for easier explanations!

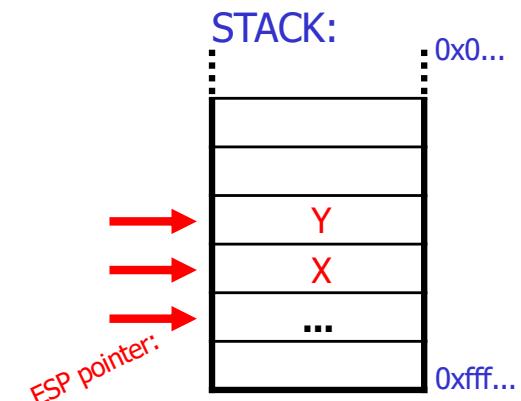


# Intel 32-bit Architecture (IA32): Basic Execution Environment

- Address space:  $1 - 2^{36}$  (64 GB),  
each process may have a linear address space of 4 GB ( $2^{32}$ )
- Basic program execution registers:
  - 8 general purpose registers (data: EAX, EBX, ECX, EDX, address: ESI, EDI, EBP, ESP)
  - 6 segment registers (CS, DS, SS, ES, FS and GS)
  - 1 flag register (EFLAGS)
  - 1 instruction pointer register (EIP)
- Stack – a continuous array of memory locations
  - Current stack is referenced by the SS register
  - ESP register – stack pointer
  - EBP register – stack frame base pointer (fixed reference)
  - PUSH – stack grows, add item (ESP decrement)
  - POP – remove item, stack shrinks (ESP increment)
- Several other registers like Control, MMX/FPU (MM/R),  
Memory Type Range Registers (MTRRs),  
SSE<sub>x</sub> (XMM), AVX (YMM), performance monitoring, ...

PUSH %eax  
PUSH %ebx  
  
<do something>  
  
POP %ebx  
POP %eax

GPRs:	
EAX:	X
EBX:	Y
ECX:	
EDX:	
ESI:	
EDI:	
EBP:	
ESP:	



# C Function Calls & Stack

- A calling function (`main`) does
  - push the parameters into stack in reverse order
  - push return address (current EIP value) onto stack
- When called, a C function (`add`) does
  - push frame pointer (EBP) into stack - saves frame pointer register and gives easy return if necessary
  - let frame pointer point at the stack top, i.e., point at the saved stack pointer (EBP = ESP)
  - shift stack pointer (ESP) upward (to lower addresses) to allocate space for local variables
- When returning, a C function (`add`) does
  - put return value in the return value register (EAX)
  - copy frame pointer into stack pointer - stack top now contains the saved frame pointer
  - pop stack into frame pointer (restore), leaving the return program pointer on top of the stack
  - the RET instruction pops the stack top into the program counter register (EIP), causing the CPU to execute from the "return address" saved earlier
- When returned to calling function, it (`main`) does
  - copy the return value (EAX) into right place
  - pop parameters – restore the stack

```
int add (int a, int b)
{
    return a + b;
}

main (void)
{
    int c = 0;
    c = add(4 , 2);
}
```



# C Function Calls & Stack

## Example:

```
int add (int a, int b)
{
    return a + b;
}
```

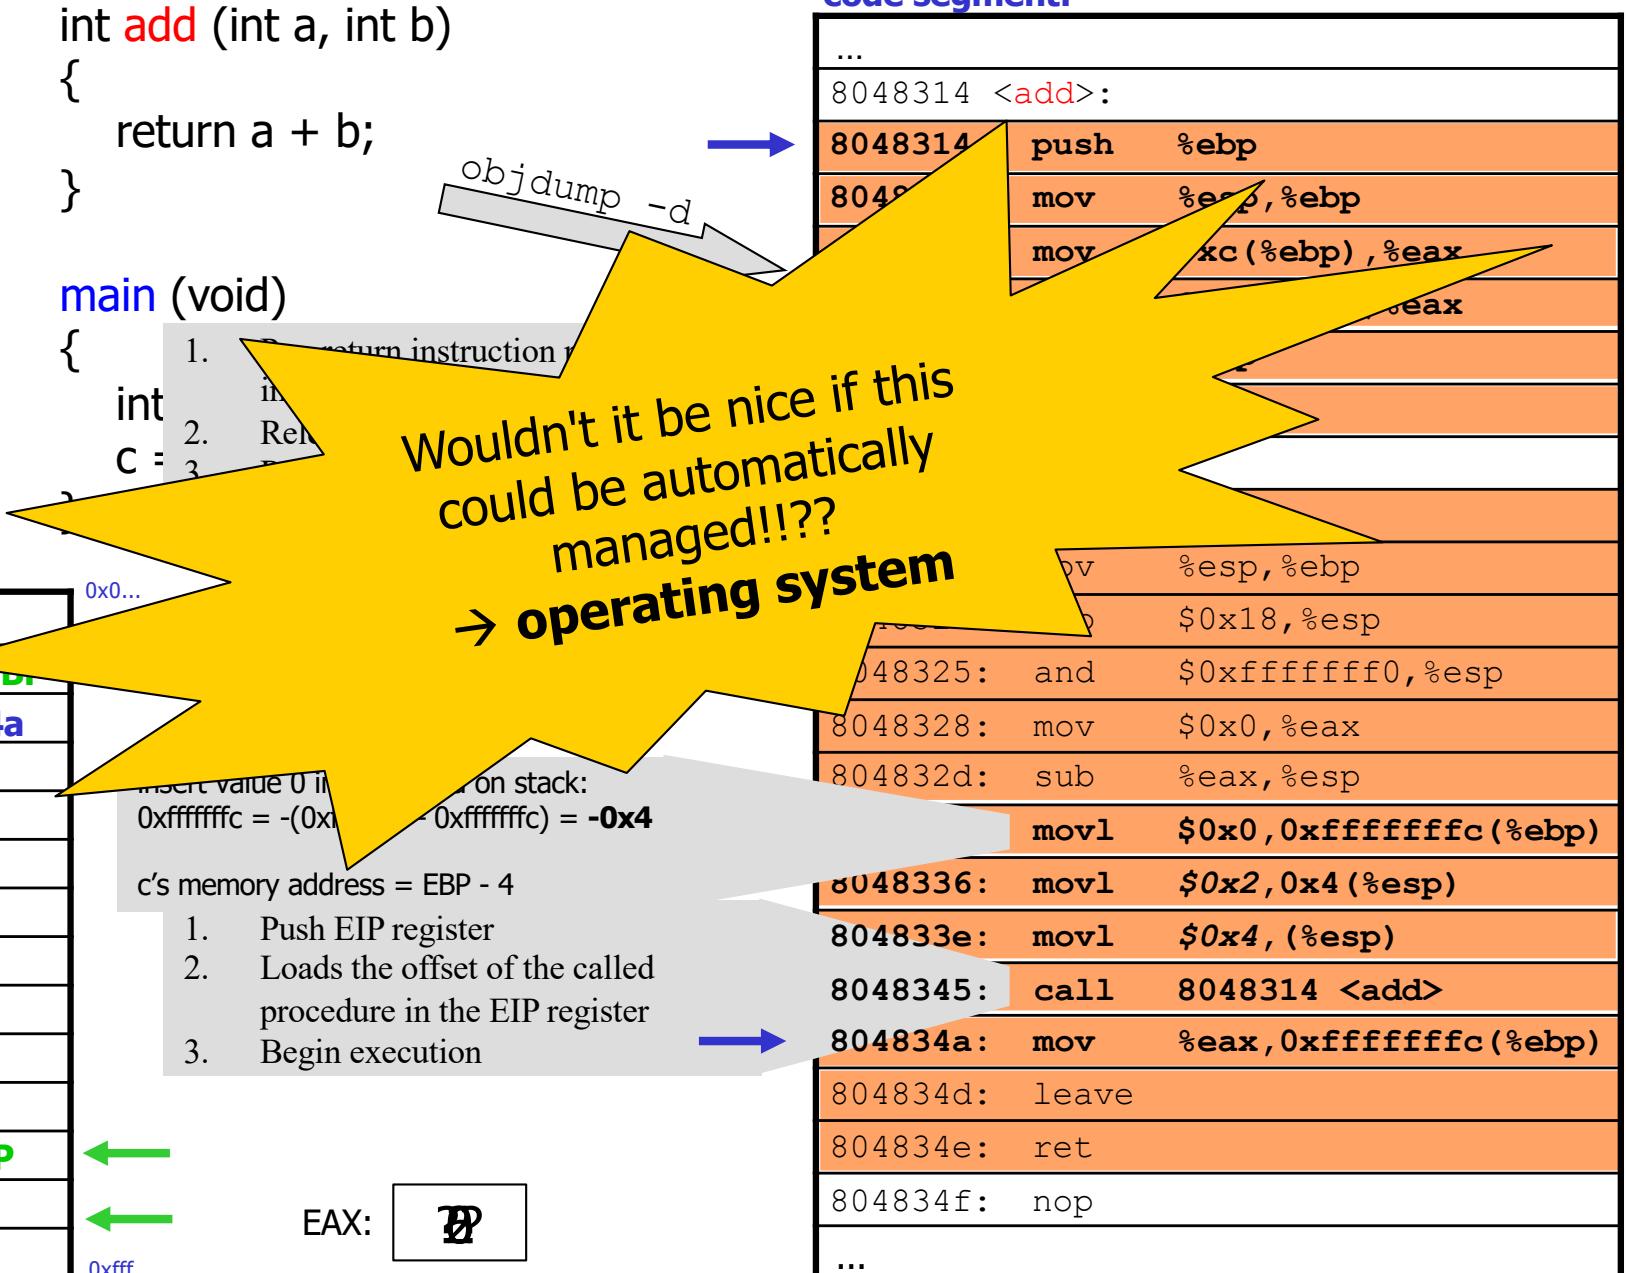
```
main (void)
```

```
{
    int
    i = 1;
    int
    c = 2;
    ...
}
```

stack:

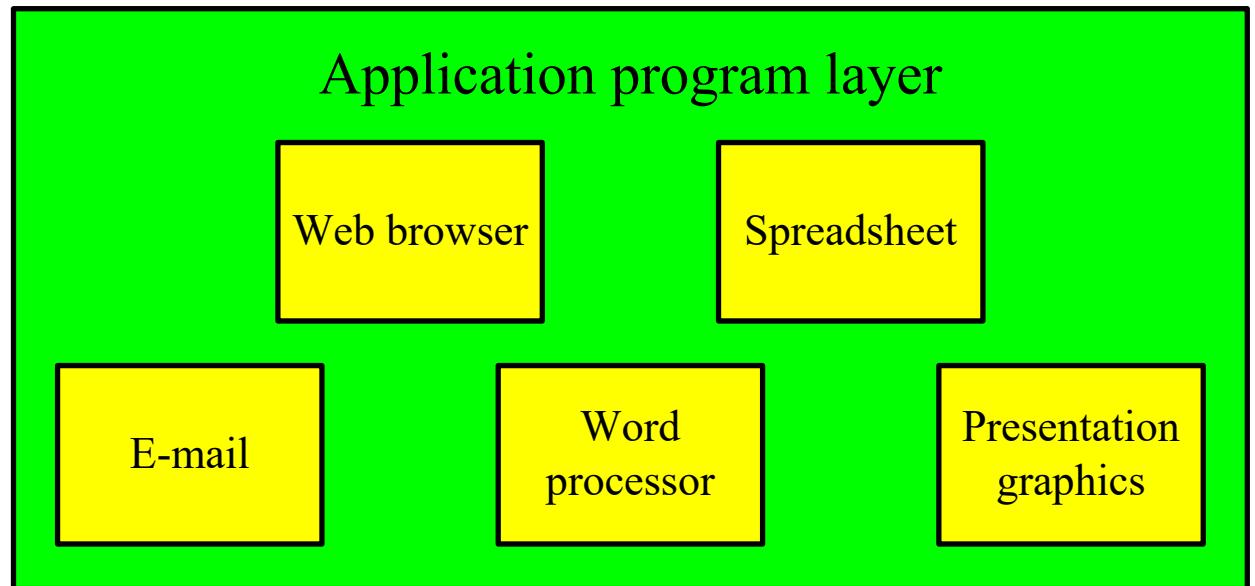
	0x0...
"main" EBP	804834a
4	4
2	2
	insert value 0 into stack: 0xfffffff0 = -(0xfffffffffc) = -0x4
	c's memory address = EBP - 4
6	6
old EBP	...
...	0xffff...

EAX: 



# Many Concurrent Tasks

- Better use & utilization
  - many concurrent processes
    - performing different tasks
    - using different parts of the machine
  - many concurrent users



- Challenges
  - “concurrent” access
  - protection/security
  - fairness
  - ...

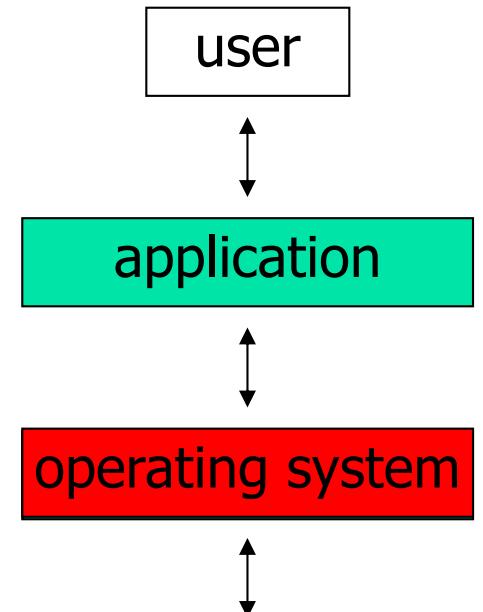


# What is an Operating System (OS)?

- "*An operating system (OS) is a collection of programs that acts as an intermediary between the hardware and its user(s), providing a high-level interface to low level hardware resources, such as the CPU, memory, and I/O devices. The operating system provides various facilities and services that make the use of the hardware convenient, efficient and safe*"

Lazowska, E. D.: Contemporary Issues in Operating Systems , in: Encyclopedia of Computer Science, Ralston, A., Reilly, E. D. (Editors), IEEE Press, 1993, pp.980

- It is an **extended machine** (top-down view)
  - Hides the messy details
  - Presents a virtual machine, easier to use
- It is a **resource manager** (bottom-up view)
  - Each program gets time/space on the resource



# Where do we find OSes?

Computers



Phones



Game Boxes



Cars



cameras,  
other vehicles/crafts,  
set-top boxes,  
watches,  
sensors,  
... → **EVERWHERE**



# Operating System Categories

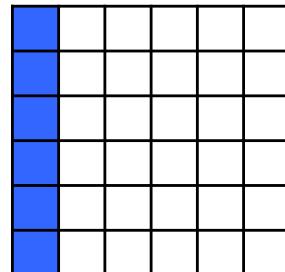
- **Single-user, single-task:**  
historic, and rare (only a few old PDAs use this)
- **Single-user, multi-tasking:**  
PCs and workstations may be configured like this (typically, phones today)
- **Multi-user, multi-tasking:**  
used on large, old mainframes; and handhelds, PCs, workstations and servers today
- **Distributed OSes:**  
support for administration of distributed resources
- **Real-time OSes:**  
support for systems with real-time requirements like cars, nuclear reactors, etc.
- **Embedded OSes:**  
built into a device to control a specific type of equipment like cellular phones, microwaves, washing machines, etc.



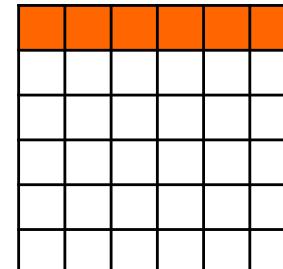
# So, why study OSes?

- “I will never write an operating system from scratch, nor even touch a line of code in the kernel”
- “Operating systems have existed for decades, what more can be added?”
- “I just need to know the API to give the system commands in order to store my data and run my programs”
- “Writing programs in Java or Python is very easy, and I do not need to know anything about operating systems to make it work”
- Consider the following example, does it matter which one to use?:

```
int data[15000][15000];
for (j = 0; j < 15000; j++)
    for (i = 0; i < 15000; i++)
        data[i][j] = i*j;
```



```
int data[15000][15000];
for (j = 0; j < 15000; j++)
    for (i = 0; i < 15000; i++)
        data[j][i] = i*j;
```



on my mac:  
datasize

3.9s	15.000	1.2s
12.2s	25.000	3.5s
22.1s	30.000	4.7s
49.3s	40.000	10.1s



# So, why study OSes?

**Failing to meet the Technical Challenges...**



... results in low quality pictures, video artifacts, hiccups, etc.

... giving **annoyed users!**



# So, why study OSes?

**Failing to meet the Technical Challenges...**



... influence the game experience

... giving **annoyed users** – latency can kill!



# So, why study OSes?

**Failing to meet the Technical Challenges...**



... influence the your everyday life

– **latency can literally kill!**

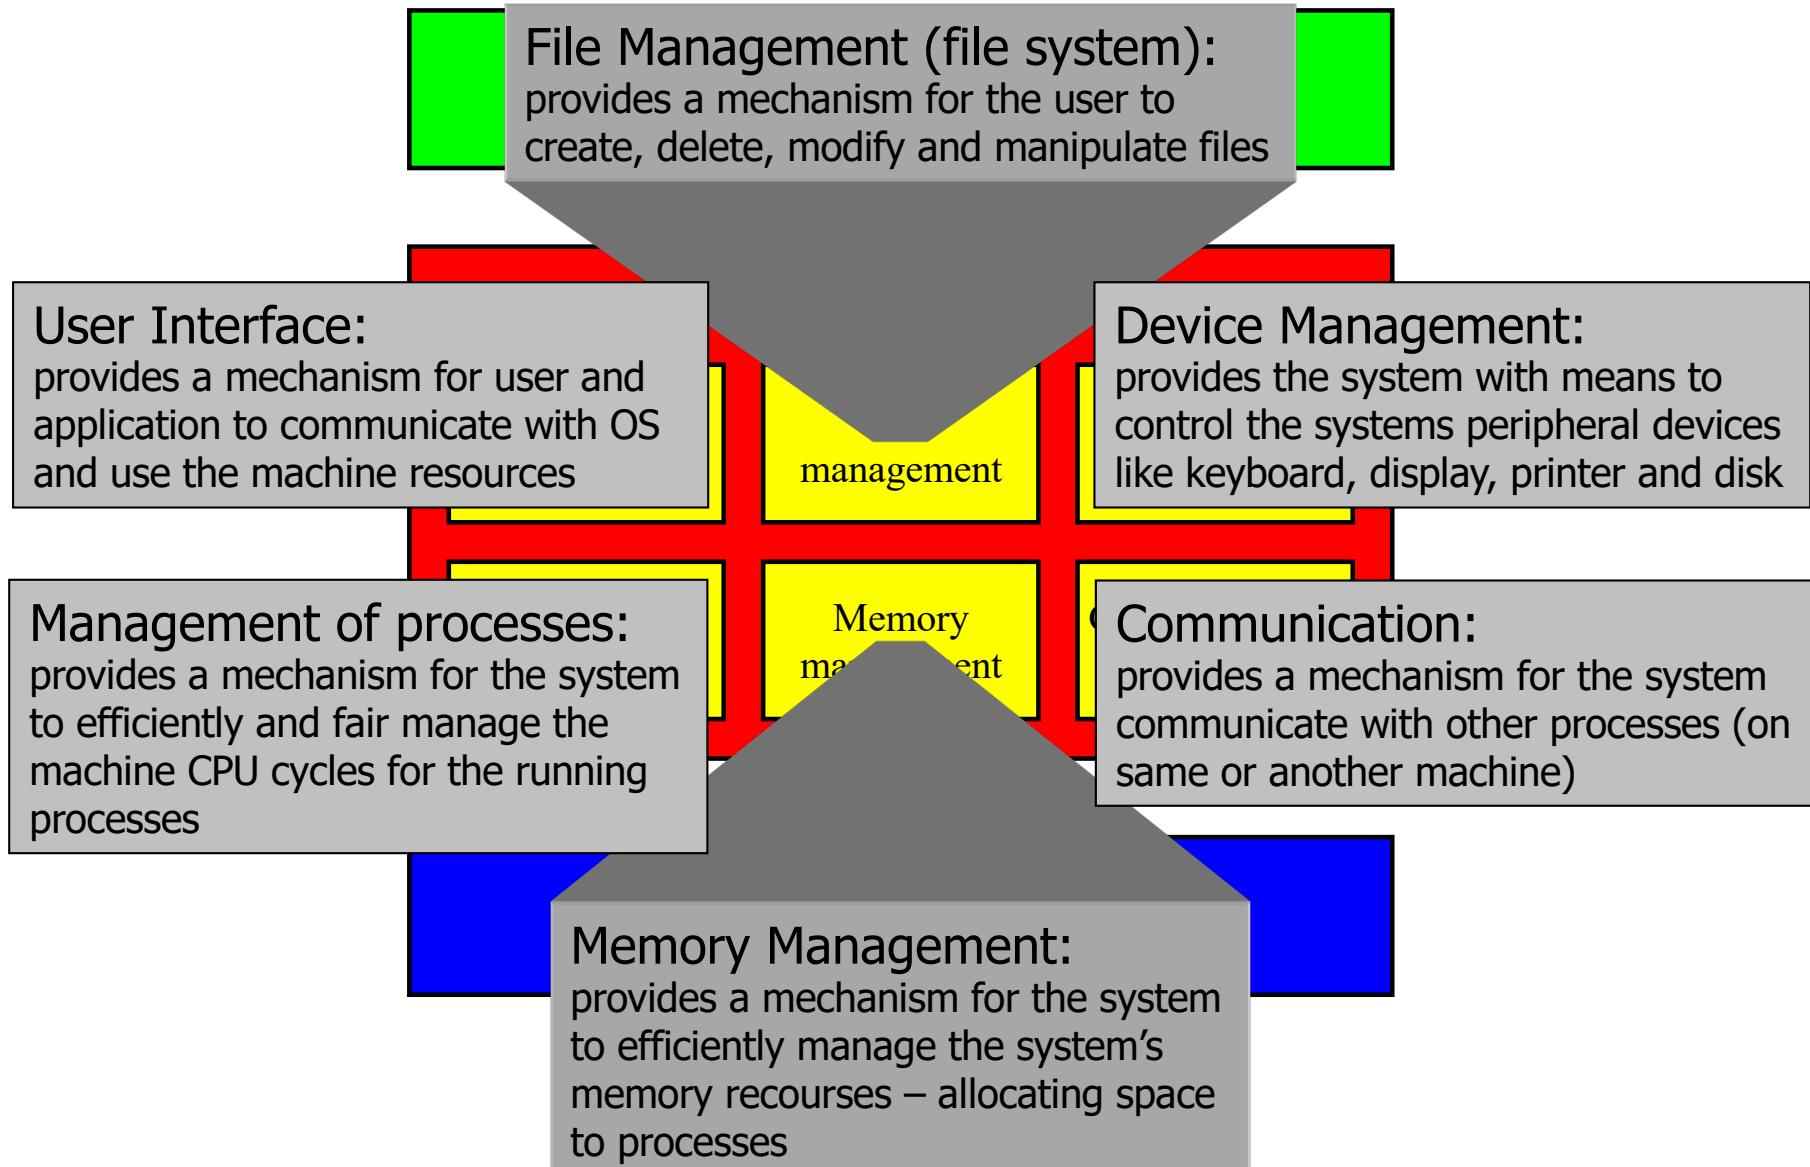


# So, why study OSes?

- To do things right and efficient, one must know how computers and operating systems work (and networks and ...)
  - operating systems provide magic to provide “infinite” CPU cycles, “endless” memory, transparent access to devices, networked computing, etc.
  - operating systems manage concurrency and sharing
  - understand the tradeoffs between performance and functionality, division of labor between HW and SW
- *OSes are found everywhere and are therefore key components in many systems*

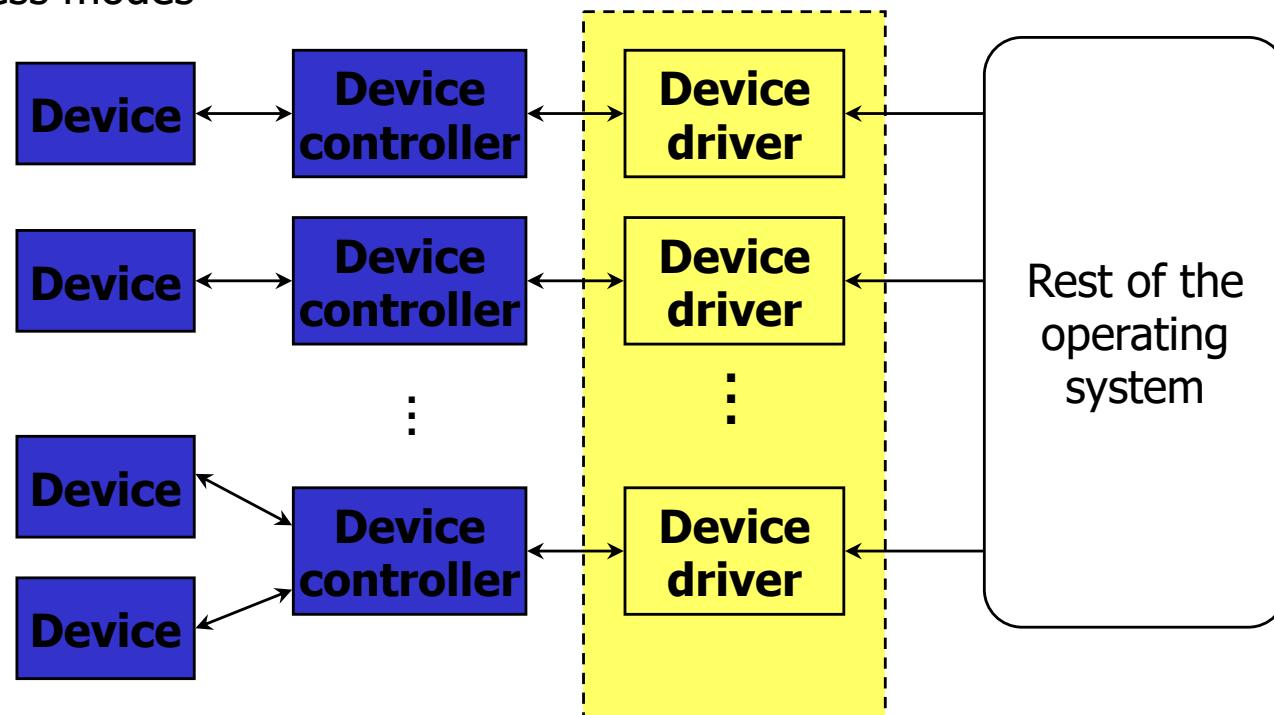


# Primary Components



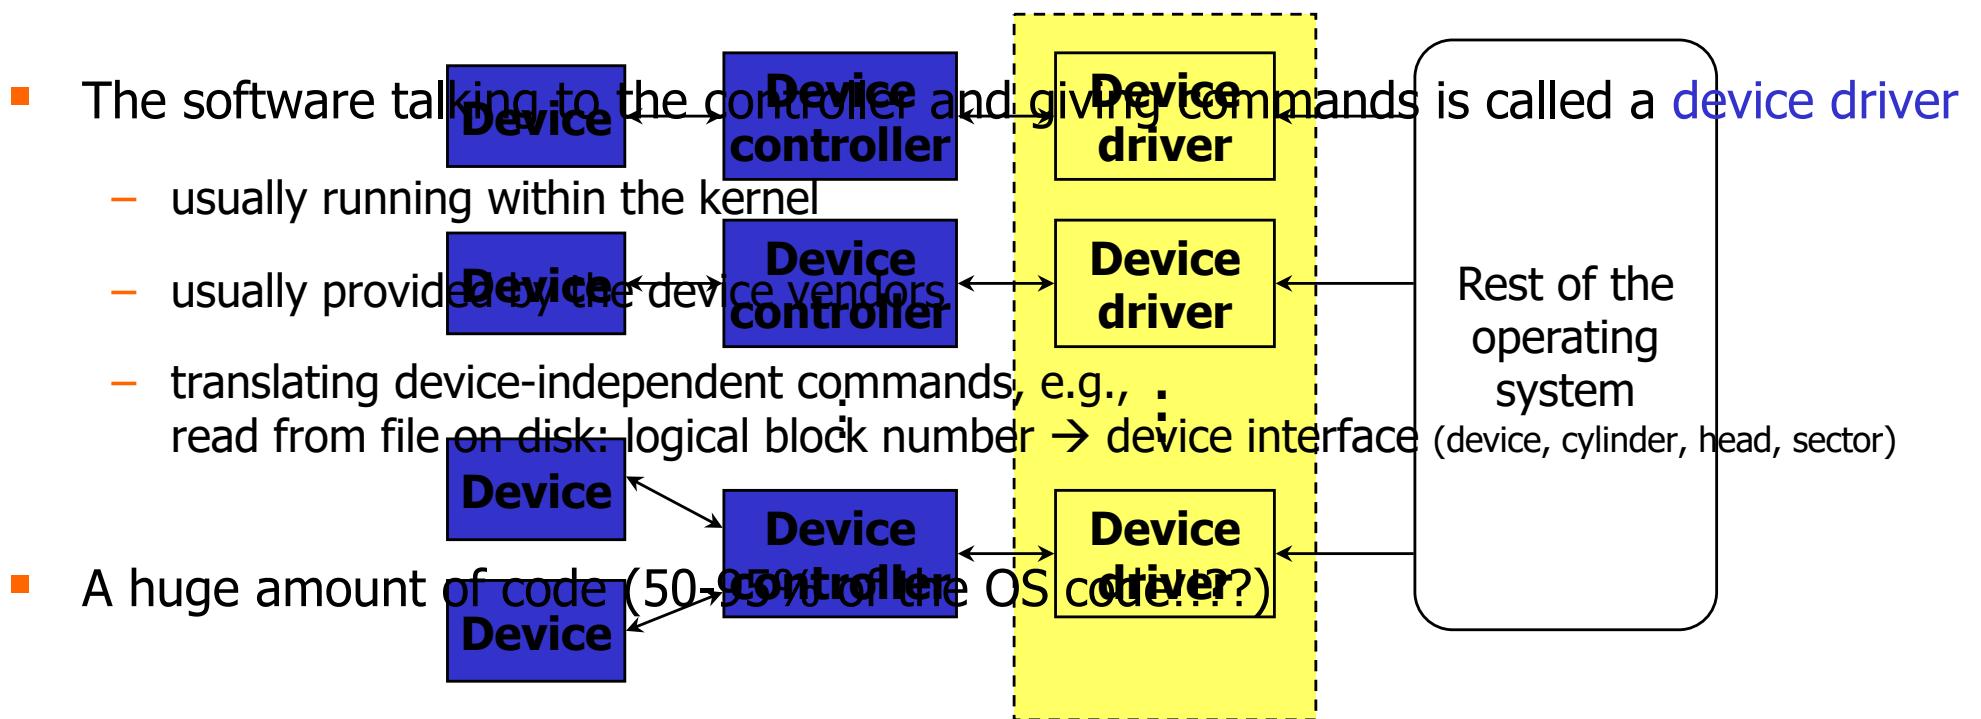
# Device Management

- The OS must be able to control peripheral devices such as disks, keyboard, network cards, screen, speakers, mouse, memory sticks, camera, DVD/Blu-Ray, microphone, printers, joysticks, ...
  - large diversity
  - varying speeds
  - different access modes



# Device Management

- Device controllers often have registers to hold status, give commands, ...
  - port I/O – special instructions to talk to device memory
  - memory mapped I/O – registers mapped into regular memory
- Each device may be different and require device-specific software



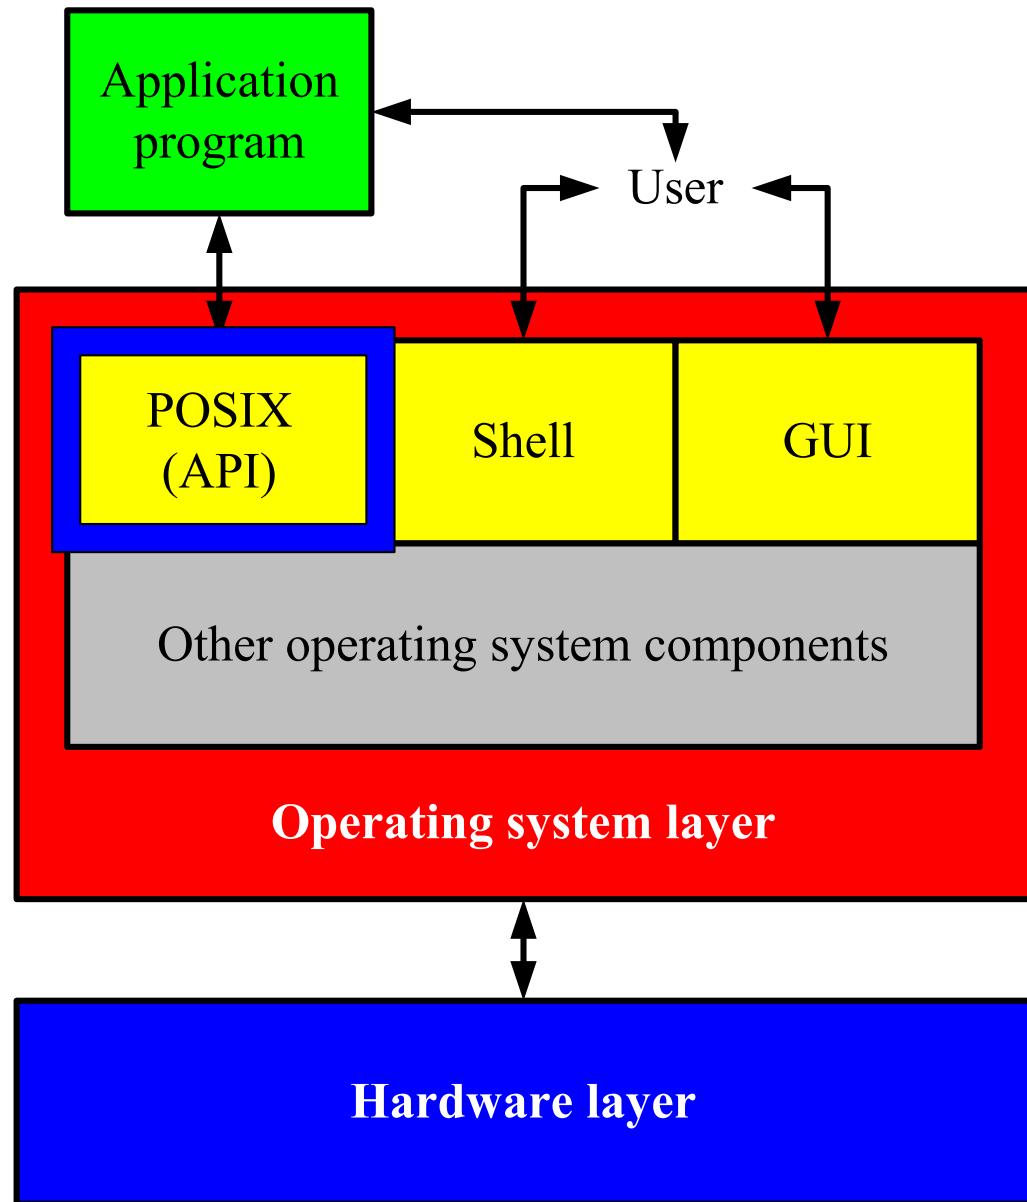
# UNIX Interfaces

Applications access HW through an API consisting of a set of routines, protocols and other tools  
(e.g., POSIX – portable OS interface for UNIX)

A user can interact with the system through the application interface or using a command line prosessed by a shell (not really a part of the OS)

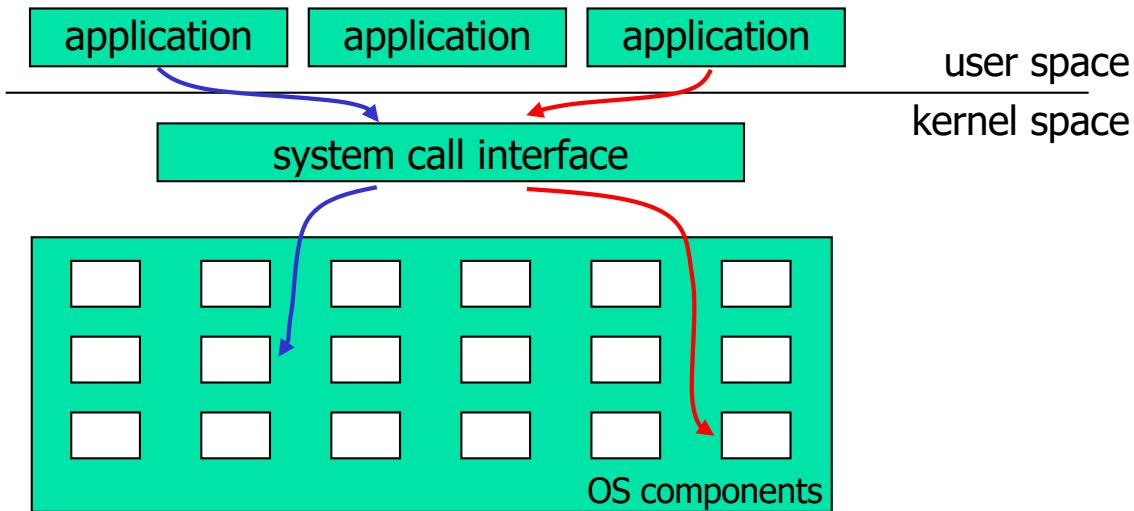
A plain command line interface may be hard to use. Many UNIX systems therefore have a standard graphical interface (X Windows) which can run a desktop system  
(like KDE, Gnome, Fvwm, Afterstep, ...)

**Windows** is more or less similar...



# System Calls

- The interface between the OS and users is defined by a set of **system calls**
- Making a system call is similar to a procedure/function call, but system calls enter the kernel:



Linux:  
x86 v2.4.19 entry.S → 242  
x86 v3.0-rc4 syscall\_table\_32.S → 347  
x86 v3.16.2 syscall\_32.tbl → 353

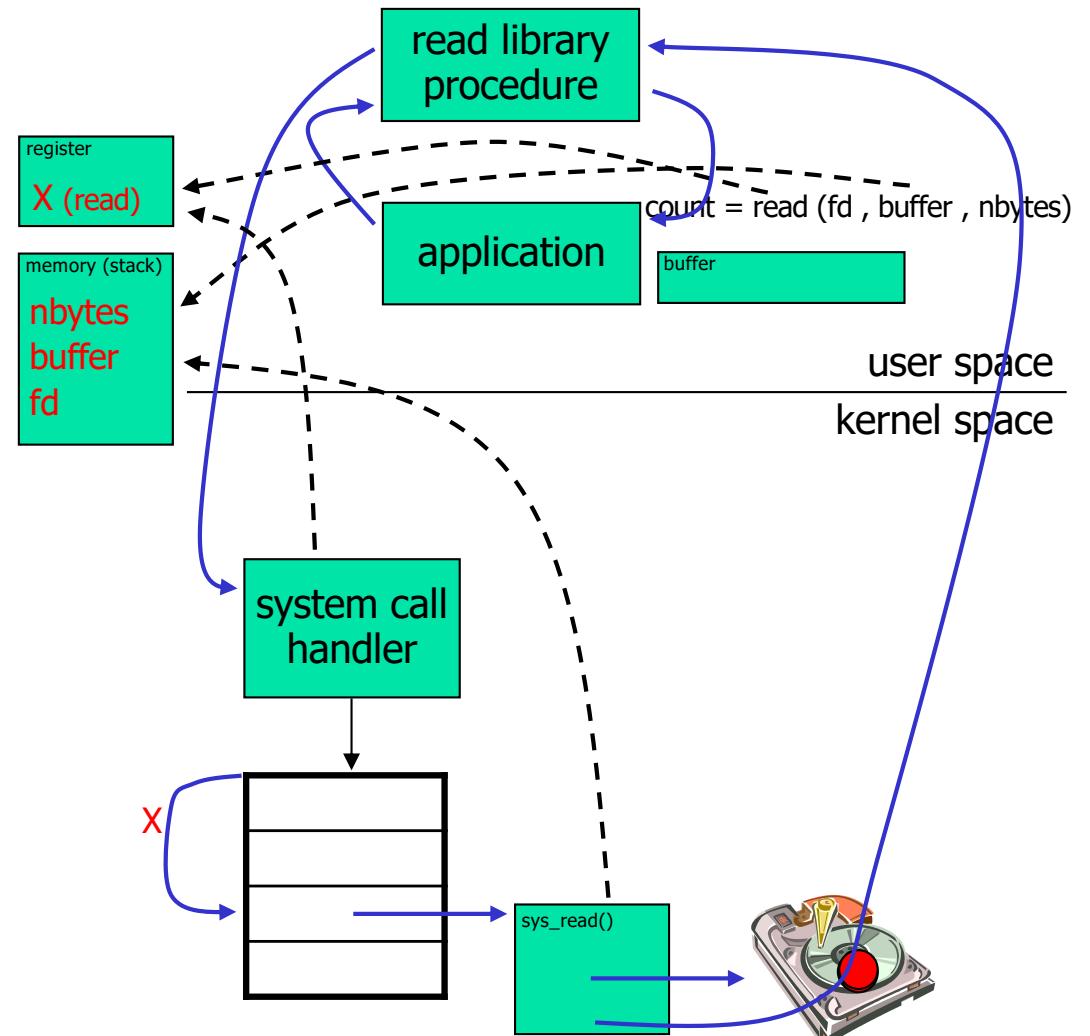
FreeBSD:  
v9 syscalls.c → 531

# System Calls: read

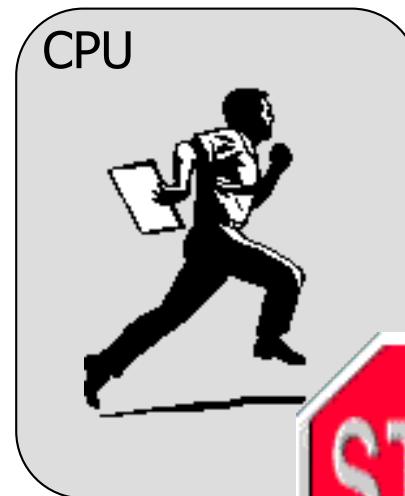
- C example:

```
count = read(fd, buffer, nbytes)
```

1. push parameters on stack
2. call library code
3. put system call number in register
4. call kernel (TRAP)
  - ✓ kernel examines system call number
  - ✓ finds requested system call handler
  - ✓ execute requested operation
5. return to library and clean up
  - ✓ increase instruction pointer
  - ✓ remove parameters from stack
6. resume process



# Interrupt Program Execution



We may use "polling"...!?

# Interrupts

---

- **Interrupts** are electronic signals that (usually) result in a forced transfer of control to an interrupt handling routine
  - alternative to polling
  - caused by *asynchronous* events like finished disk operations, incoming network packets, expired timers, ...
  - an interrupt descriptor table (IDT) associates each interrupt with a code descriptor (pointer to code segment)
  - can be disabled or masked out



# Exceptions

---

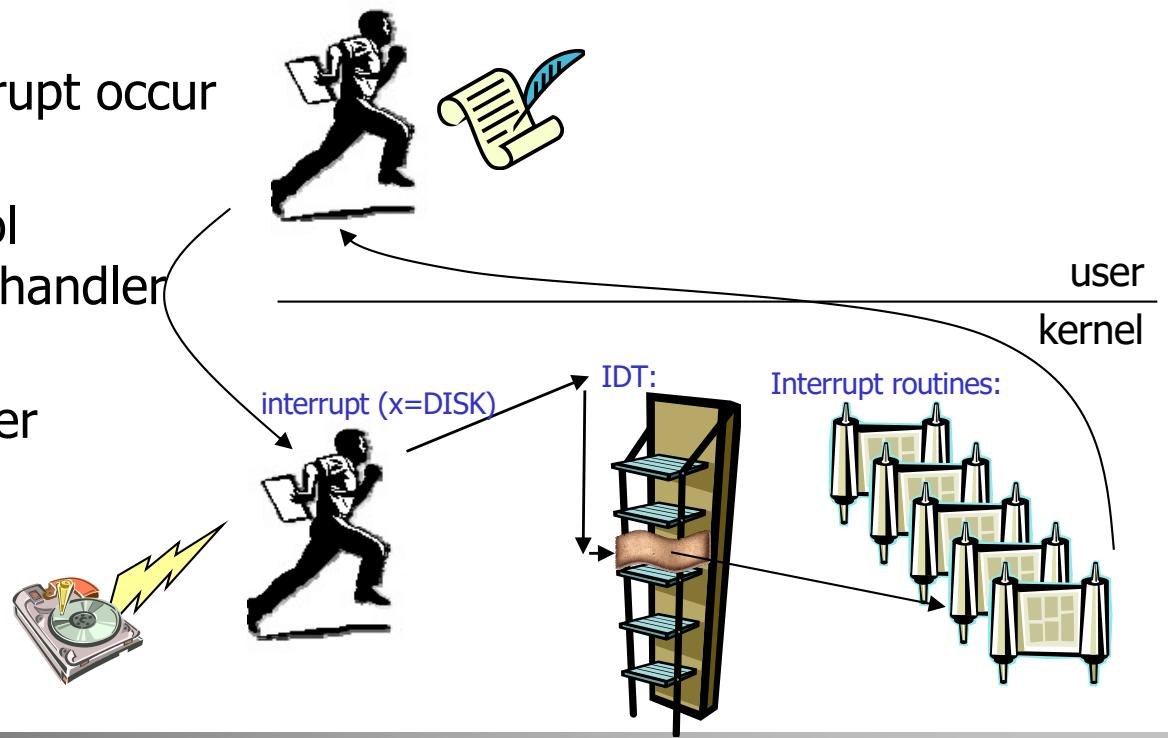
- Another way for the processor to interrupt program execution is **exceptions**
  - caused by *synchronous* events generated when the processor detects a predefined condition while executing an instruction
  - **TRAPS**: the processor reaches a condition the exception handler can handle (e.g., overflow, break point in code like making a system call, ...)
  - **FAULTS**: the processor reaches a fault the exception handler can correct (e.g., division by zero, wrong data format, ...)
  - **ABORTS**: terminate the process due to an unrecoverable error (e.g., hardware failure) which the process itself cannot correct
  - the processor responds to exceptions (i.e., traps and faults) essentially as for interrupts



# Interrupt (and Exception) Handling

- The IA-32 has an interrupt description table (IDT) with 256 entries for interrupts and exceptions
  - 32 (0 - 31) predefined and reserved
  - 224 (32 - 255) is "user" (operating system) defined
- Similar to system calls, each interrupt is associated with a code segment through the IDT and a unique index value giving management like this:

- process running while interrupt occur
- capture state, switch control and find the right interrupt handler
- execute the interrupt handler
- restore interrupted process
- continue execution



# Booting

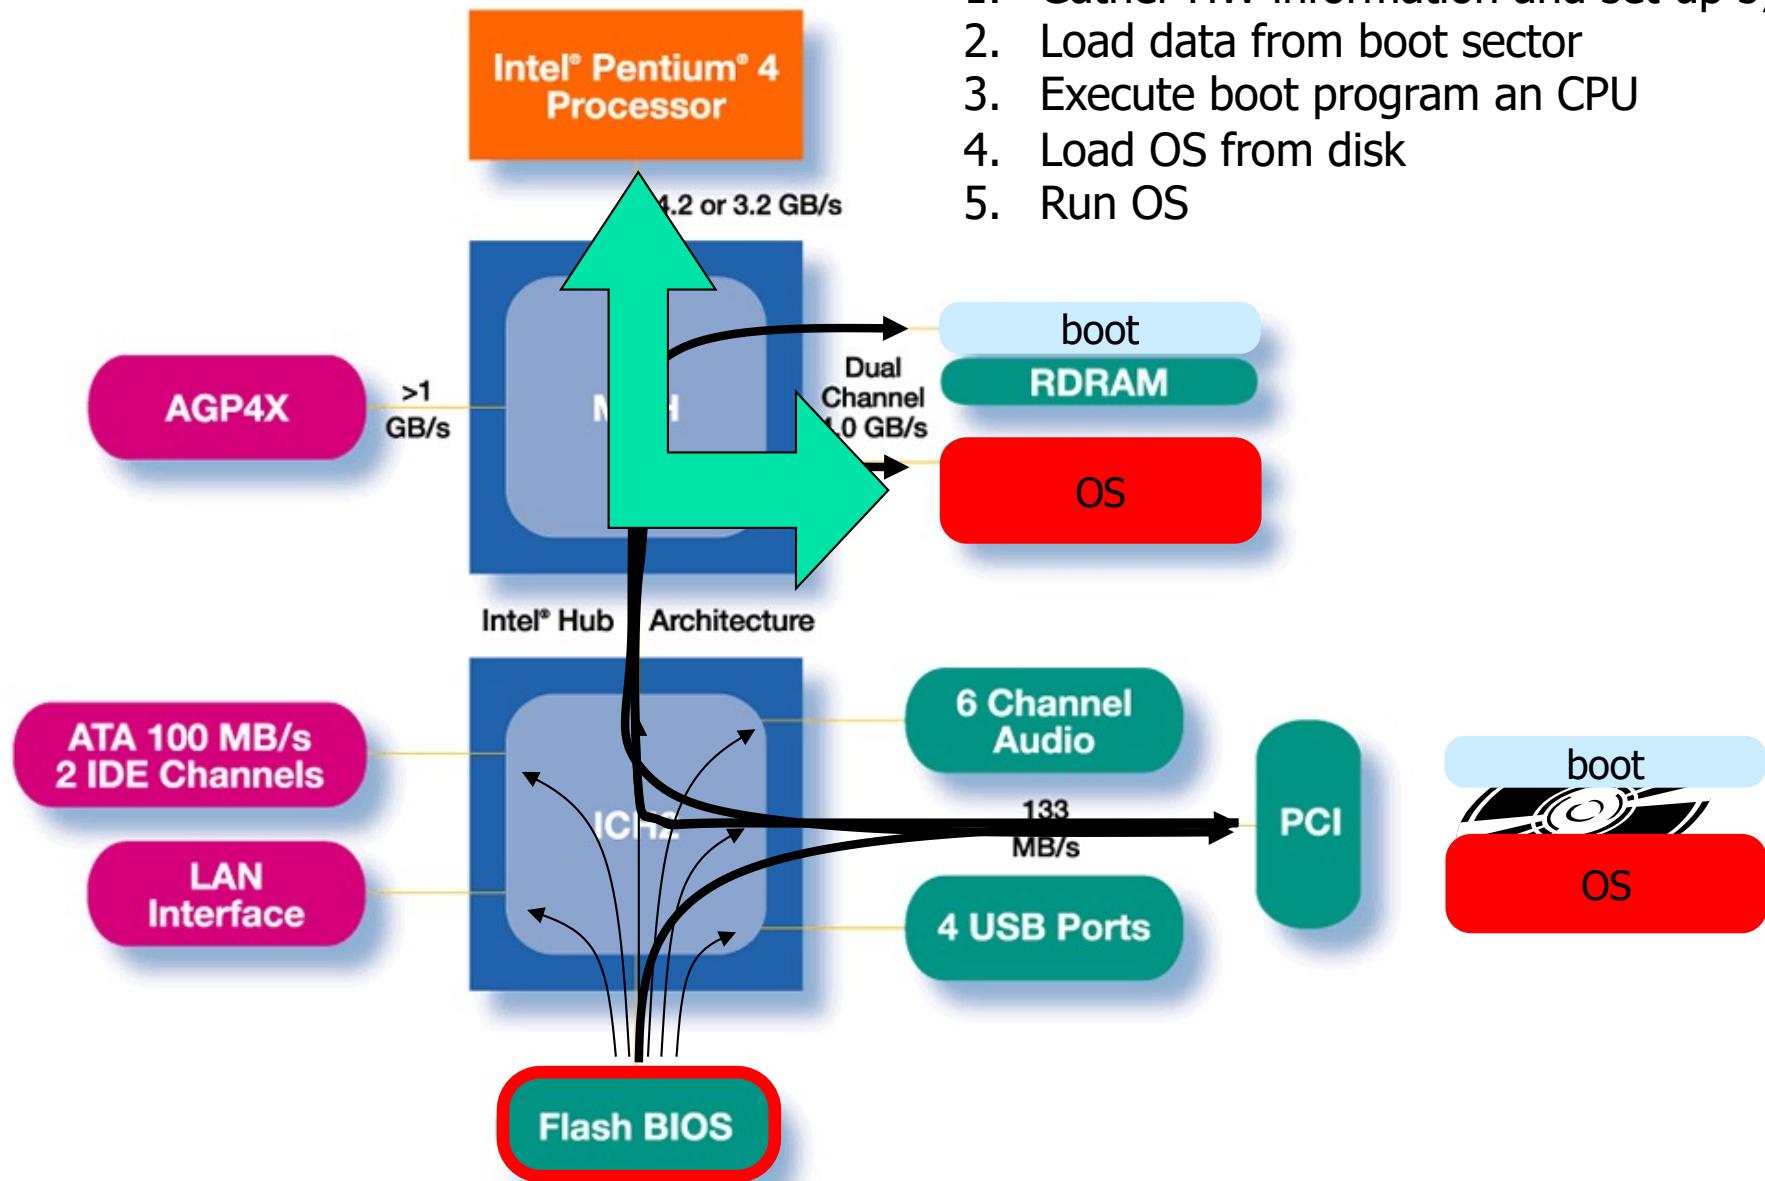
---

- Memory is a volatile, limited resource: OS usually on disk
- Most motherboards contain a **basic input/output system** (BIOS) chip (often flash RAM) – stores instructions for basic HW initialization and management, and initiates the ...
- ... **bootstrap**: loads the OS into memory
  - read the **boot** program from a known location on secondary storage typically first sector(s), often called **master boot record** (MBR)
  - run **boot** program
    - read root file system and locate file with OS kernel
    - load kernel into memory
    - run kernel



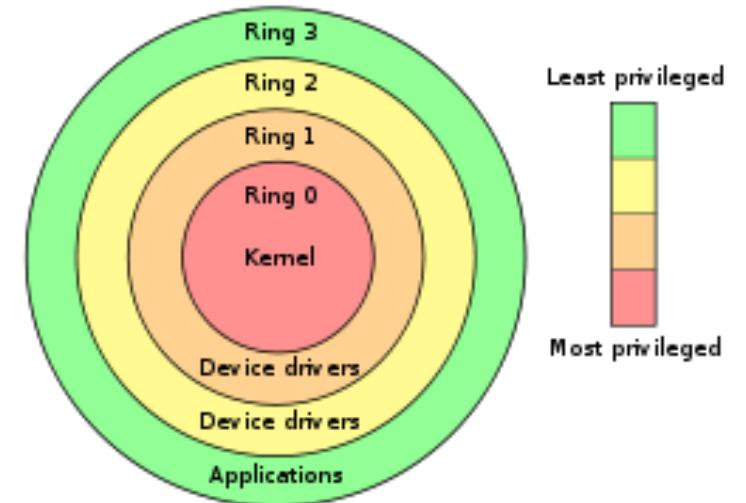
# Booting

1. Gather HW information and set up system
2. Load data from boot sector
3. Execute boot program an CPU
4. Load OS from disk
5. Run OS



# User Level vs. Kernel Level (Protection)

- Many OSes distinguish user and kernel level, i.e., due to security and protection
- Usually, applications and many sub-systems run in user mode (level 3)
  - protected mode
  - not allowed to access HW or device drivers directly, only through an API
  - access to assigned memory only
  - limited instruction set
- OSes run in kernel mode (under the virtual machine abstraction, level 0)
  - real mode
  - access to the entire memory
  - all instructions can be executed
  - bypass security

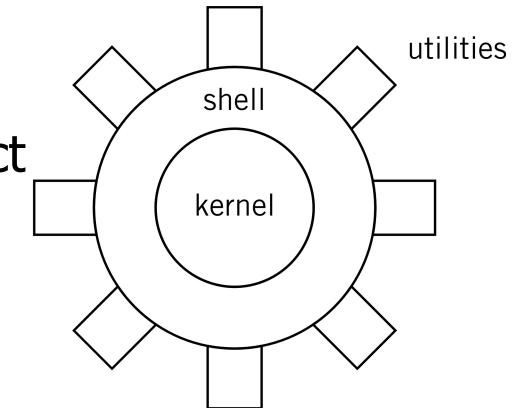


# OS Organization

- No standard describing how to organize a kernel (as it is for compilers, communication protocols, etc.) and several approaches exist, e.g.:

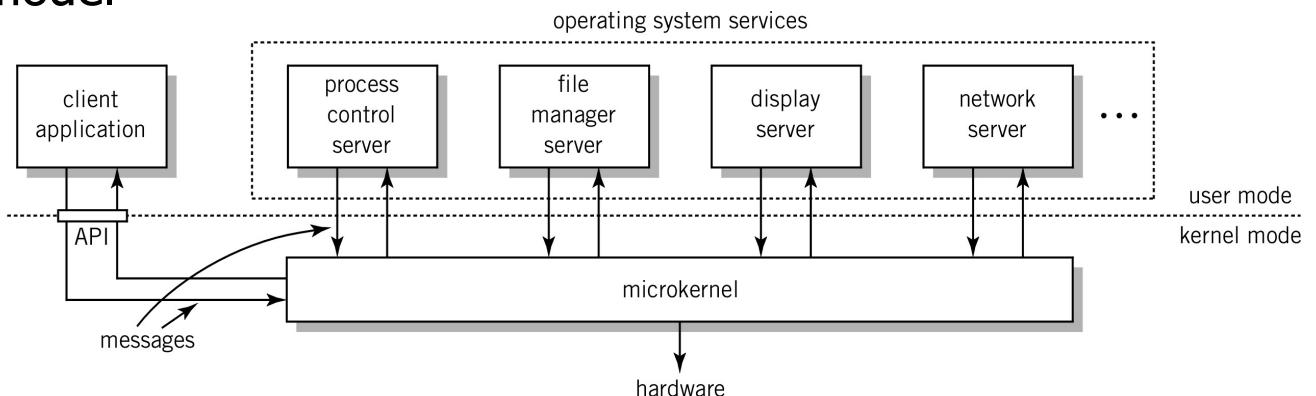
- **Monolithic kernels** ("the big mess"):

- written as a collection of functions linked into a single object
- usually efficient (no boundaries to cross)
- large, complex, easy to crash
- UNIX, Linux, Windows 7 (++), ...



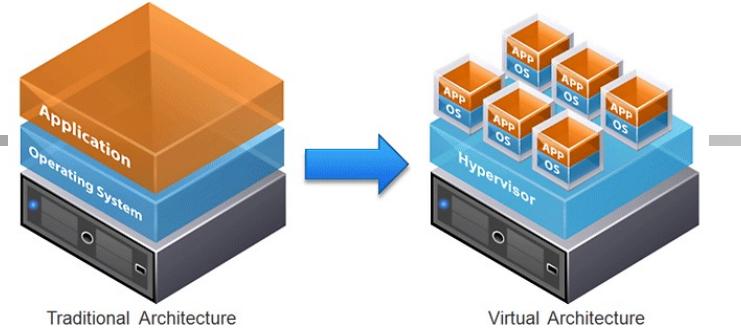
- **Micro kernels**

- kernel with minimal functionality (managing interrupts, memory, processor)
- other services are implemented in server processes running in user space used in a client-server model
- lot of message passing (inefficient)
- small, modular, extensible, portable, ...
- MACH, L4, Chorus, Windows NT, ...



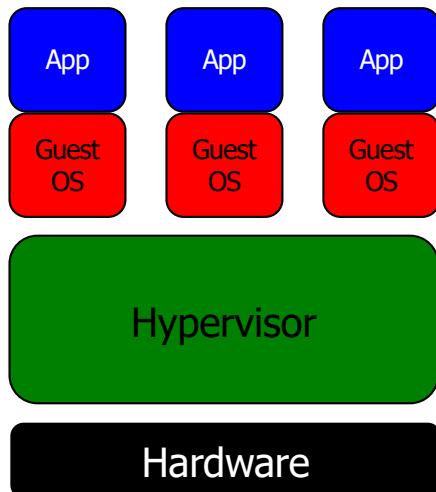
# Virtualization

- People would like to save money, save energy, reduce the number of machines, be secure, easily move services, etc.... and still run multiple OSes, applications, etc...

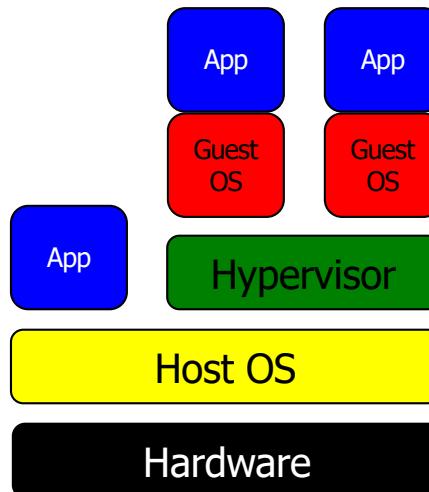


## Virtualization

- many types of virtualization – server/machine virtualization
- partitioning a physical server into several virtual servers, or machines
- interact independently/isolated with other devices, applications, data and users

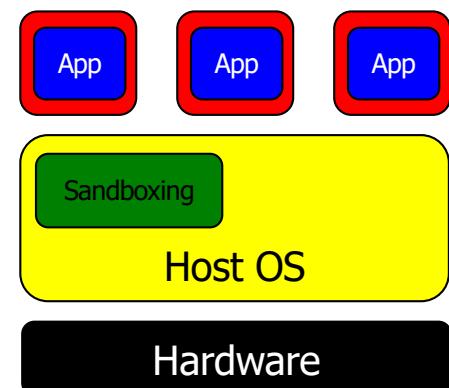


**Type 1 Hypervisor**  
VMWare ESX, XEN, Hyper-V



**Type 2 Hypervisor**  
KVM, VirtualBox, VMWare Workstation

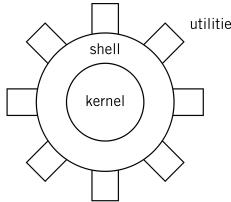
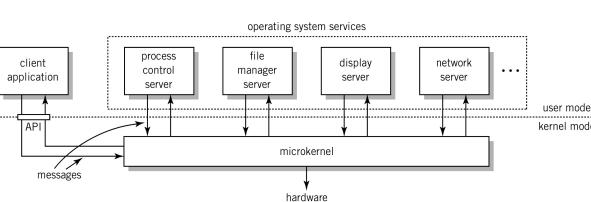
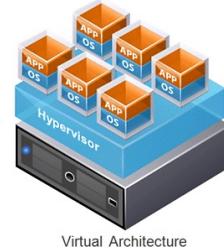
## Sandboxing



**Sandboxing**  
sandboxie, bufferzone, libcontainer, runc, cgroups



# System Structures

Monolithic kernel	Micro kernels	Virtual Machines
		
Performance Easy to share resources	Flexible, modular Failure isolation	Flexibility Many virtual computers - consolidation Isolation
Unstructured – big mess <b>One failure kills it all</b> (everything is in real (privileged) mode)	Performance ? – multiple boundary crossings (but claim switching and IPC just cost some 10s-100s cycles)	Performance ? - multiple boundary crossings Complexity ?



# Summary

- OSes are found “everywhere” and provide **virtual machines** and work as a **resource managers**
- Many components provide different services
- Users access the services using an interface like system calls
- In the next lectures, we look closer at some of the main components and abstractions in an OS
  - processes management
  - memory management
  - storage management
  - local inter-process communication
  - inter-computer network communication is covered in the last part of the course

